


Predicting Customer Lifetime Value (CLV) Using Python


Data-Driven Approach to Estimate Future Customer Spending

Presented By:
2023BCS109 Pruthviraj Jadhav
2023BCS101 Krishna Kalyan


Under the Guidance of:
Prof. Amit Nandedkar Sir



Agenda




Introduction



Primary Goals



Plan



Output



Summary

What is CLV

Customer Lifetime Value (CLV) is a metric that estimates the total net profit a company can generate from a customer over their entire relationship.

Understanding CLV allows you to make informed decisions based on how long a customer typically buys from you and what they spend over the lifetime of that relationship. This metric can help inform your strategy on acquisition, customer retention, customer support, and even the quality of your products and services.



Customer Lifetime Value: The Metric You Can't Afford to Ignore

“ “

Eighty percent of our business
comes from 20% of our customers.
It costs 10 times less to sell to an
existing customer than to find a new
customer

” ”

Why CLV_



Eighty percent of our business comes from 20% of our customers.

It costs 10 times less to sell to an existing customer than to find a new customer

Businesses prefer retaining customers over acquiring new ones.

Businesses are curious to weigh the value of each customer and determine who is truly worth investing in this is exactly when we use customer lifetime value to identify the total worth of a customer based on their relationship.

Primary goal

Predict a customer's total spending in the next year based on their first few purchases.



Packages Used

Numpy
high-speed
mathematical
operations

Scikitlearn
Most imp tool for
model training

Matplotlib
Turns data into clear
visual stories


Pandas
for data cleaning,
analysis with
DataFrames

Openpyxl
Reads, edits, and
creates Excel files
programmatically


Plan Connecting Dots




In the search of
Find a most fitting dataset to perform analysis



Taking useful out of it
Data Cleaning & Analysis of cleaned data



Playing with Data
Model Training using Linear Regression & Random Forest Algorithms




Final Output
Output results, visual analysis of predictions.



Future Scope
Customer Segmentation & Power BI Dashboards

Found a perfect match



Data Set Overview

Sources:

Online Retail Dataset (UCI ML Repository)
E-commerce Transactions Data (Kaggle)

It is a transnational data set which contains all the transactions occurring between 01/12/2010 and 09/12/2011 for a UK-based and registered non-store online retail. The company mainly sells unique all-occasion gifts. Many customers of the company are wholesalers.

It contains 8 columns:

InvoiceNo, StockCode, Description, Quantity, InvoiceDate, UnitPrice, CustomerID, Country

Data Cleaning & Analysis of cleaned data

Imported raw transactional data (data.csv) with 541,909 records.

```
>>> df = pd.read_csv("data.csv", encoding="latin1")
>>> print(df.head())
   InvoiceNo StockCode  ... CustomerID      Country
0      536365    85123A  ...     17850.0  United Kingdom
1      536365     71053  ...     17850.0  United Kingdom
2      536365    84406B  ...     17850.0  United Kingdom
3      536365    84029G  ...     17850.0  United Kingdom
4      536365    84029E  ...     17850.0  United Kingdom

[5 rows x 8 columns]
>>> df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
 #   Column        Non-Null Count  Dtype  
--- 
 0   InvoiceNo     541909 non-null   object 
 1   StockCode     541909 non-null   object 
 2   Description   540455 non-null   object 
 3   Quantity      541909 non-null   int64  
 4   InvoiceDate   541909 non-null   object 
 5   UnitPrice     541909 non-null   float64
 6   CustomerID    406829 non-null   float64
 7   Country       541909 non-null   object 
dtypes: float64(2), int64(1), object(5)
memory usage: 33.1+ MB
```

Data Cleaning & Analysis of cleaned data

```
df.dropna(subset=["CustomerID"], inplace=True)  
df["InvoiceDate"] = pd.to_datetime(df["InvoiceDate"])  
df["TotalSales"] = df["Quantity"] * df["UnitPrice"]
```

Removed missing CustomerID values: removed rows with missing CustomerID

Converted timestamps: Converted **InvoiceDate** into proper datetime format for time-based analysis.

Created TotalSales column: Created a new feature TotalSales=Quantity × UnitPrice for revenue per transaction.

Data Cleaning & Analysis of cleaned data

Final dataset: 406,829 valid transactions across 8 features + 1 derived feature.

```
>>> df['InvoiceDate']=pd.to_datetime(df['InvoiceDate'])
>>> print(df['InvoiceDate'].dtypes)
datetime64[ns]
>>> df['TotalSales']=df['Quantity']*df['UnitPrice']
>>> df.info()
<class 'pandas.core.frame.DataFrame'>
Index: 406829 entries, 0 to 541908
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   InvoiceNo    406829 non-null   object 
 1   StockCode    406829 non-null   object 
 2   Description  406829 non-null   object 
 3   Quantity     406829 non-null   int64  
 4   InvoiceDate  406829 non-null   datetime64[ns]
 5   UnitPrice    406829 non-null   float64
 6   CustomerID   406829 non-null   float64
 7   Country      406829 non-null   object 
 8   TotalSales   406829 non-null   float64
dtypes: datetime64[ns](1), float64(3), int64(1), object(4)
memory usage: 31.0+ MB
>>>
```

How we get there

simple and effective methodology used in calculating customer value over a time frame is RFM_

Recency

R

How recently a customer has made a purchase

Frequency

F

How often a customer makes a purchase

Monetary

M

Dollar value of the purchases

calculating

Recency

Frequency

Monetary

```
def calculate_rfm(df):
    """Calculate Recency, Frequency, and Monetary metrics."""
    snapshot_date = df["InvoiceDate"].max() + pd.Timedelta(days=1)
    rfm = df.groupby("CustomerID").agg({
        "InvoiceDate": lambda x: (snapshot_date - x.max()).days,
        "InvoiceNo": "nunique",
        "TotalSales": "sum"
    })
```

snapshot date : Sets reference date as day after last transaction in dataset

Groups all rows by CustomerID

For each customer group, finds their latest purchase date


Calculates days between snapshot and last purchase

Counts number of unique invoice numbers per customer

Adds up all Total Sales values per customer.


Model Training

LINEAR REGRESSION



Linear Regression
Model 1

Random Forest



Random Forest
Model 2

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
```

Make essential imports


Then, we divide data in 2 parts_
2/3 for training
1/3 for testing

Models learns the relationship between features & target from 2/3 rd part and predicts the output which is later compared with actual values from 1/3rd part so as to calculate the accuracy.

The math is mathing_

Math inside it


LINEAR REGRESSION



```
linreg = LinearRegression().fit(X_train, y_train)
```

$$\text{FutureMonetary} = a \times \text{Recency} + b \times \text{Frequency} + c \times \text{Monetary} + \text{intercept}$$

Random Forest



```
rf = RandomForestRegressor(n_estimators=100, random_state=42).fit(X_train, y_train)
```

Creates 100 decision trees and combines them

$$\hat{y} = \beta_0 + \beta_1 R + \beta_2 F + \beta_3 M ; \beta_1 = ? , \beta_2 = ? ; \beta_3 = ?$$

To solve this, we need to solve -

$$\beta = (X^T X)^{-1} X^T y$$

where β = coeff. vector

y = target vector (actual future Monetary from test dataset)

$X = [\beta_0 \ R \ F \ M]$ for each customer
 ↴ constant for all (let = 1)

The linear model : $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2$
 ∴ in matrix form: $\hat{y} = X \beta$

$$[\]_{1 \times 1} = [\beta_0 \ R \ F \ M]_{1 \times 4} * \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \beta_3 \end{bmatrix}_{4 \times 1}$$

To minimize sum of squared Errors -

$$\text{cost function: } J(\beta) = \sum (y_i - \hat{y}_i)^2 = (y - X\beta)^T (y - X\beta) \\ = (y^T - \beta^T X^T) (y - X\beta) \\ = y^T y - 2\beta^T X^T y + \beta^T X^T X \beta$$

$$\frac{\partial J(\beta)}{\partial \beta} = \frac{\partial}{\partial \beta} (y^T y - 2\beta^T X^T y + \beta^T X^T X \beta) = 0$$

$$\Rightarrow -2 X^T y + 2 X^T X \beta = 0$$

solve for $\beta =$

$$X^T y = X^T X \beta \\ \therefore \beta = (X^T X)^{-1} X^T y$$

Example: suppose after training we get

$$\left. \begin{array}{l} \beta_0 = 50.0 \\ \beta_1 = -2.0 \\ \beta_2 = 25.0 \\ \beta_3 = 0.8 \end{array} \right\} \text{predicted formula -}$$

$$\therefore \hat{y} = 50.0 + (-2.0)R + (25.0)F + (0.8)M$$

|r

n -estimators = 100 \Rightarrow creates 100 decision trees

\Rightarrow Random Forest creates 100 different training subsets:

Tree 1 : [0, 1, 3, 4, 5, 7]


Tree 2 : [2, 3, 4, 5, 6, 7] customers

Tree 3 : [0, 1, 2, 5, 6, 7]

⋮

Consider Tree 1 = ALGORITHM tries all possible splits across all features -

complete Tree 1 structure -



first split on basis of Recency

then on Frequency

then on Monetary.

Tree 2, 3 ... 100 splits vary. splits will also change the order




Hence, final prediction = $\left(\frac{1}{100}\right) * \sum \text{Tree Prediction}$


rf


Final outputs

```
plot_monthly_sales,  
plot_top_products,  
plot_top_countries,  
plot_predictions(y_test, y_pred_lr, "Linear Regression", "blue")  
plot_predictions(y_test, y_pred_rf, "Random Forest", "green")
```









Future scope

Create an **interactive dashboard (Tableau/Power BI)** for marketing managers to visualize CLV predictions in real-time.

Use unsupervised learning (K-Means) on the RFM features to identify **5-7 distinct customer segments** for highly targeted marketing.





Thank you.